Accuracy sees the evidence; confidence does not

Paired within each of 373 questions and token-matched — same duration, same frame count, same presented and live token counts. The two windows differ only in whether the evidence is inside. The accuracy effect varies 7.0× across the bins; the confidence effect follows none of that variation and crosses zero at every one of them.

A Removing the evidence opens a gap in accuracy, and almost none in confidence

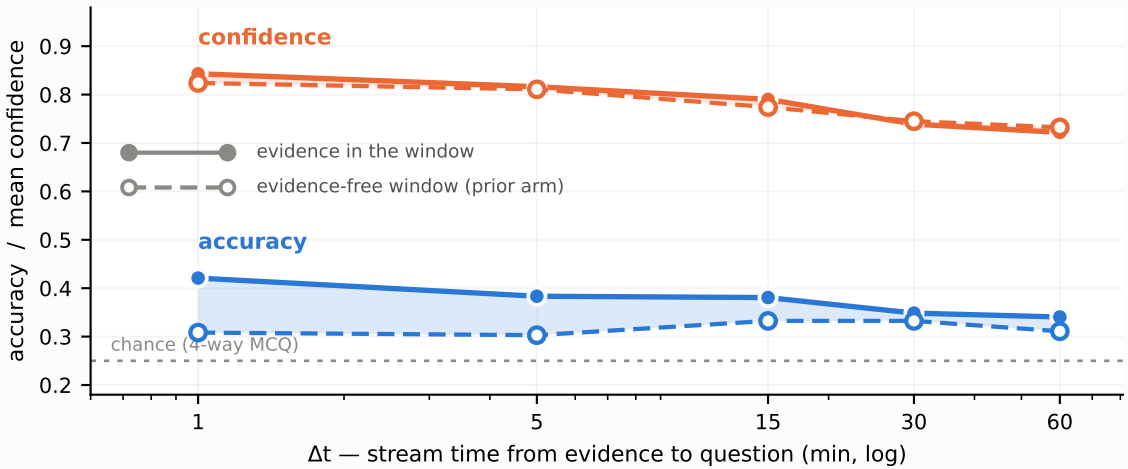

B Filled = interval excludes zero. Every accuracy mark is filled; no confidence mark is

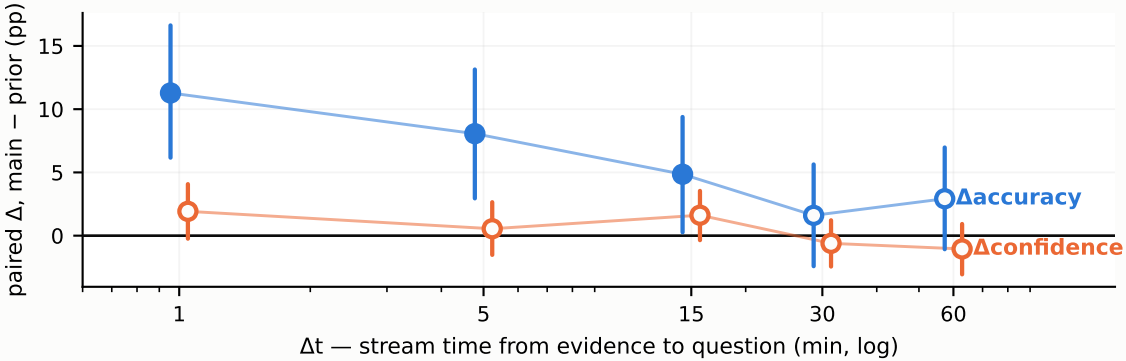

Confidence distributions overlap 0.88 / 0.83 / 0.84 / 0.85 / 0.90 across the bins, and confidence separates the two conditions at AUROC 0.511 [0.495, 0.528] on 1865 pairs, against 0.500 for chance. The probe is not inert: the same measure moves +12.6 [+10.8, +14.4] on question kind. Confidence is the option-letter softmax, not a calibrated probability.

Drawn by egovm/streaming/x12_evidence_figure.py from x11_evidence_blindness.json; no number is computed here.