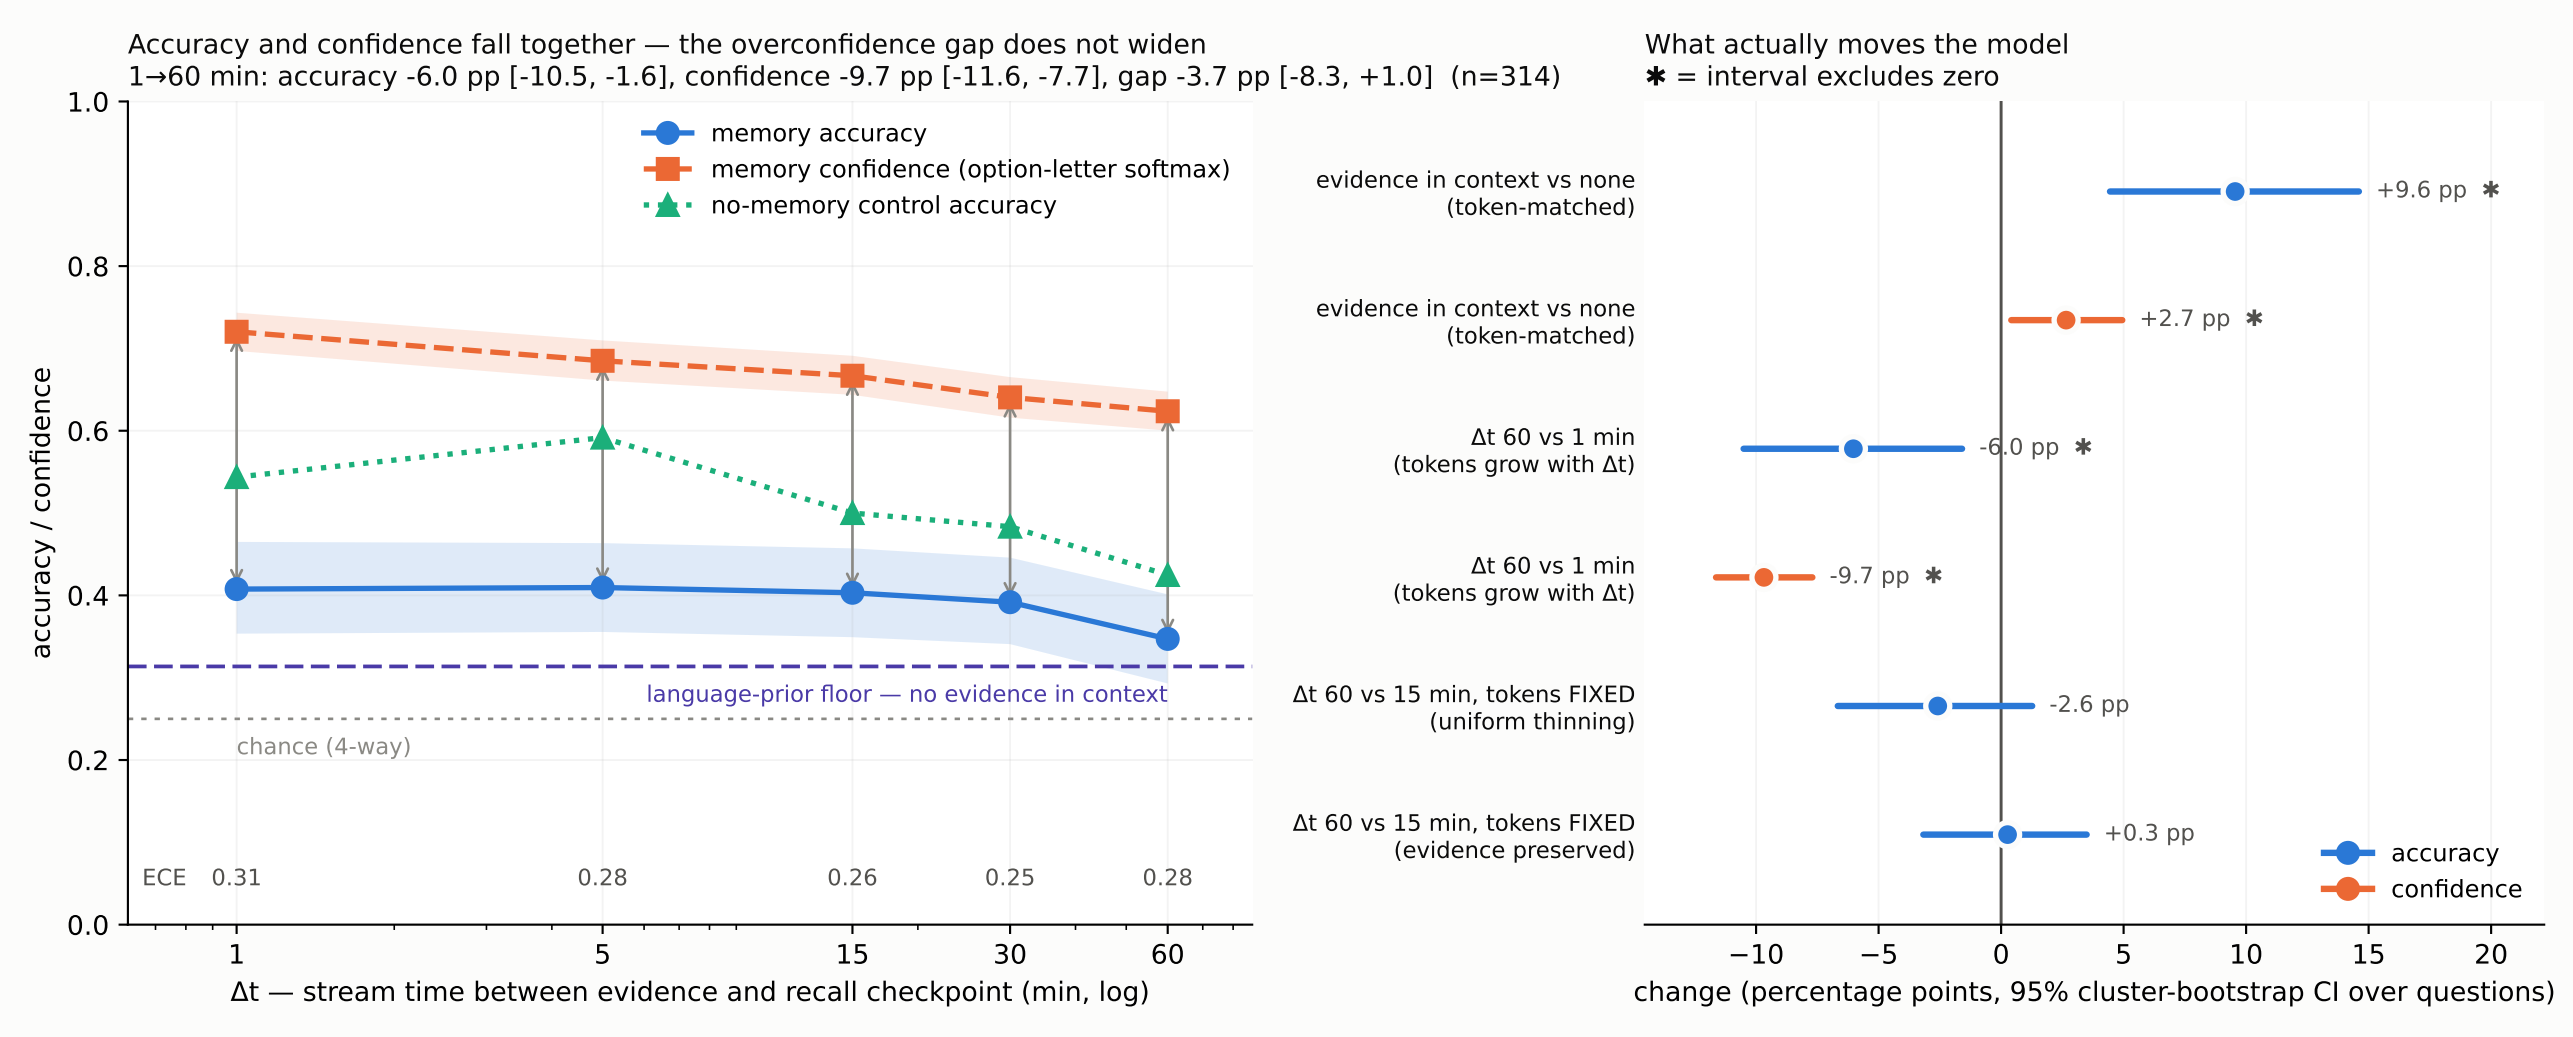
Accuracy and confidence fall together — the overconfidence gap does not widen
1→60 min: accuracy -6.0 pp [-10.5, -1.6], confidence -9.7 pp [-11.6, -7.7], gap -3.7 pp [-8.3, +1.0] (n=314)
memory accuracy
memory confidence (option-letter softmax)
no-memory control accuracy
language-prior floor — no evidence in context
chance (4-way)
ECE 0.31 0.28 0.26 0.25 0.28
accuracy / confidence
Δt — stream time between evidence and recall checkpoint (min, log)
What actually moves the model
✱ = interval excludes zero
evidence in context vs none (token-matched)
+9.6 pp ✱
evidence in context vs none (token-matched)
+2.7 pp ✱
Δt 60 vs 1 min (tokens grow with Δt)
-6.0 pp ✱
Δt 60 vs 1 min (tokens grow with Δt)
-9.7 pp ✱
Δt 60 vs 15 min, tokens FIXED (uniform thinning)
-2.6 pp
Δt 60 vs 15 min, tokens FIXED (evidence preserved)
+0.3 pp
accuracy
confidence
change (percentage points, 95% cluster-bootstrap CI over questions)