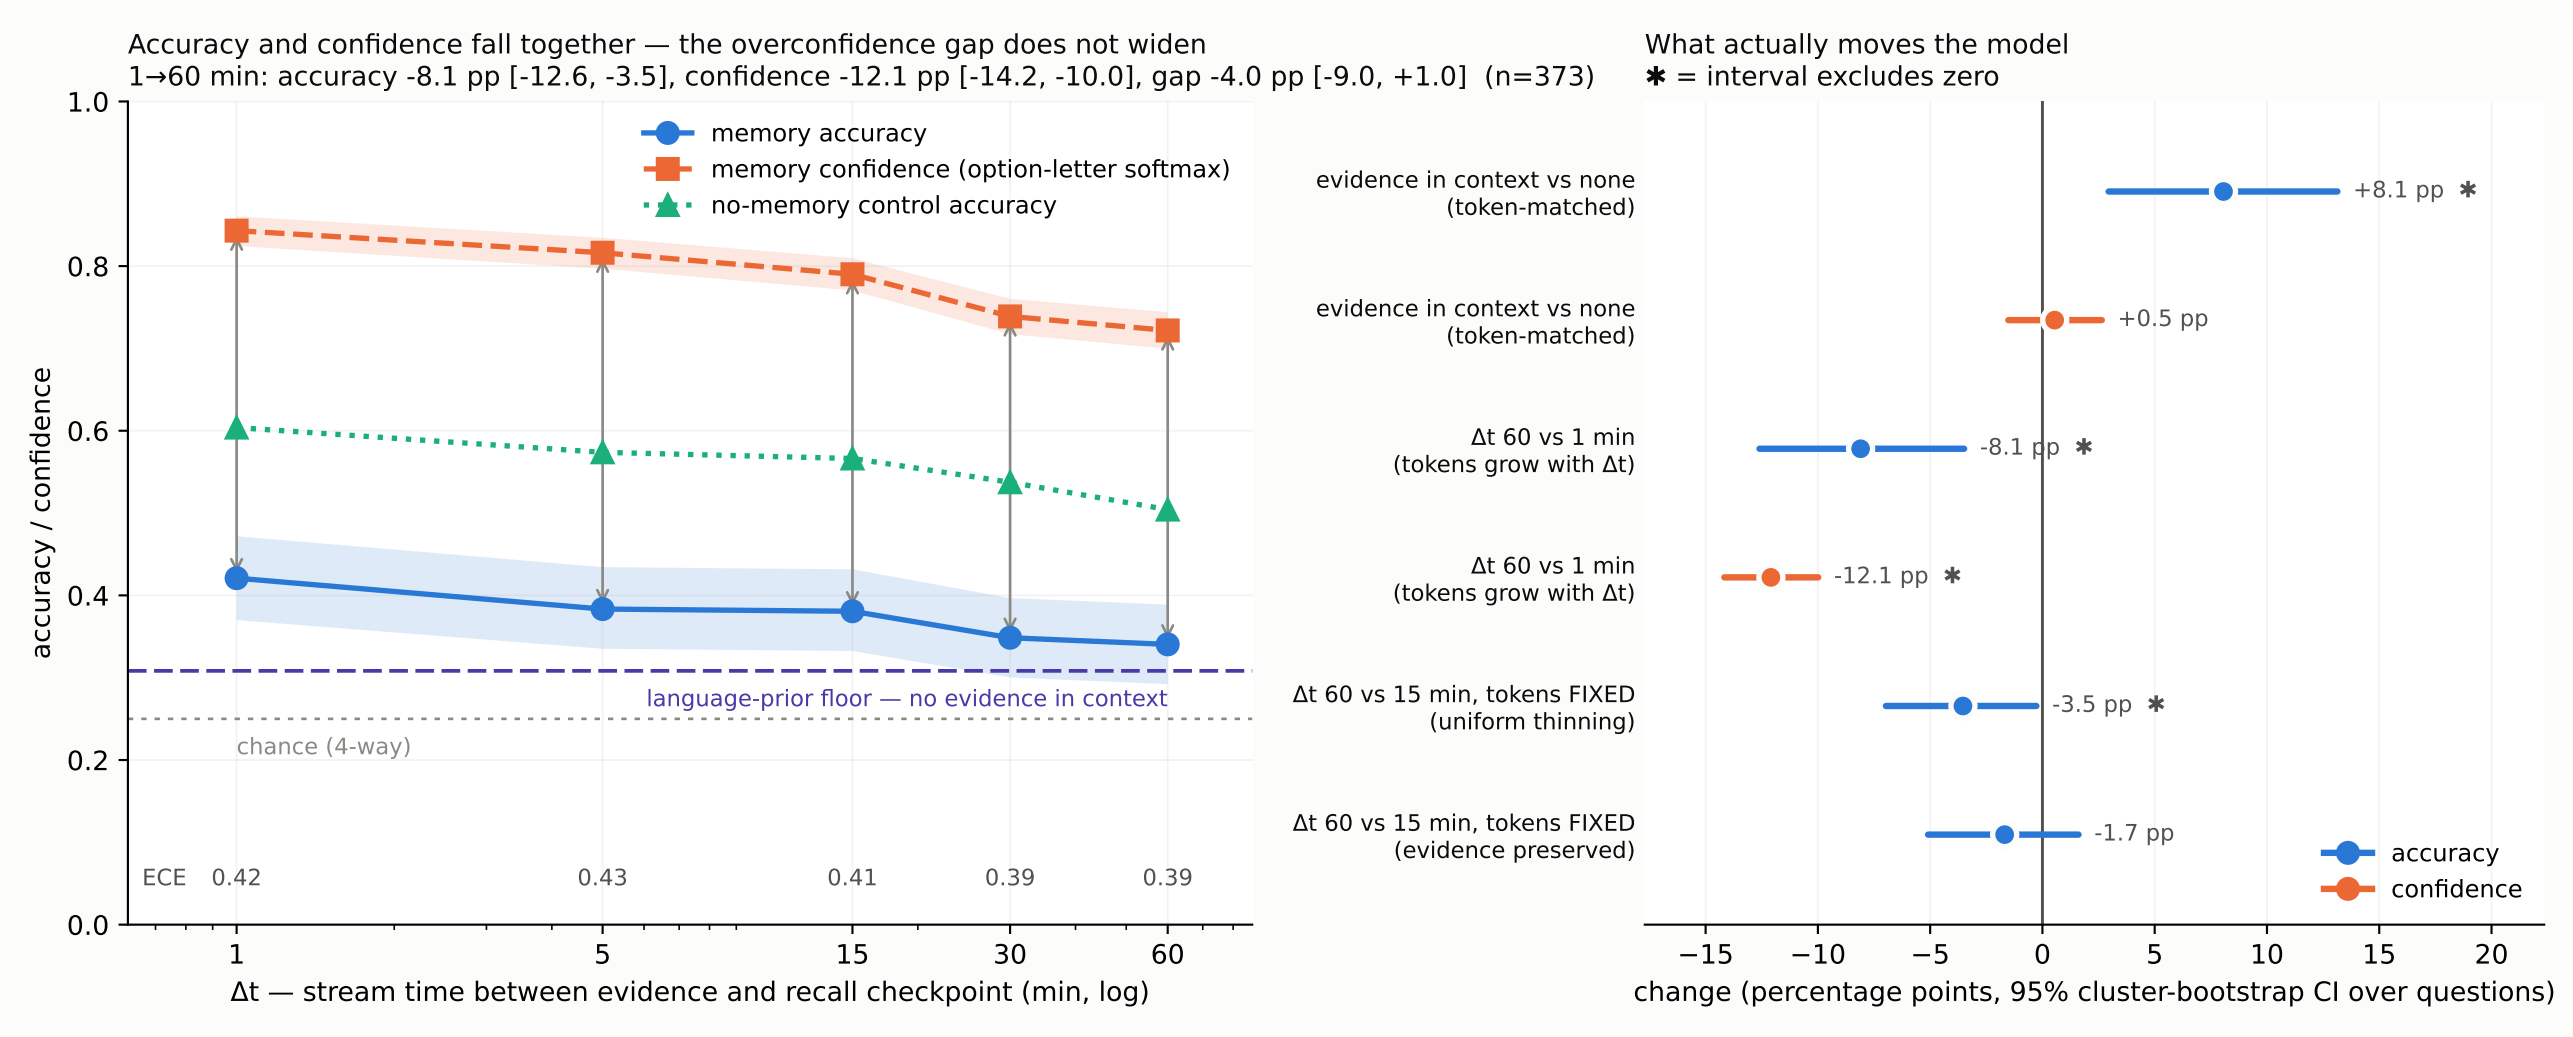
Accuracy and confidence fall together — the overconfidence gap does not widen
1→60 min: accuracy -8.1 pp [-12.6, -3.5], confidence -12.1 pp [-14.2, -10.0], gap -4.0 pp [-9.0, +1.0] (n=373)
memory accuracy
memory confidence (option-letter softmax)
no-memory control accuracy
language-prior floor — no evidence in context
chance (4-way)
ECE 0.42 0.43 0.41 0.39 0.39
accuracy / confidence
Δt — stream time between evidence and recall checkpoint (min, log)
What actually moves the model
✱ = interval excludes zero
evidence in context vs none (token-matched)
+8.1 pp ✱
evidence in context vs none (token-matched)
+0.5 pp
Δt 60 vs 1 min (tokens grow with Δt)
-8.1 pp ✱
Δt 60 vs 1 min (tokens grow with Δt)
-12.1 pp ✱
Δt 60 vs 15 min, tokens FIXED (uniform thinning)
-3.5 pp ✱
Δt 60 vs 15 min, tokens FIXED (evidence preserved)
-1.7 pp
accuracy
confidence
change (percentage points, 95% cluster-bootstrap CI over questions)